

## BIU Slurm System

**Slurm** is a Linux open-source resource manager and job scheduler designed to utilize a fair share of the computing resources.

### What Are Slurm Partitions?

**Partitions** in Slurm are sets of compute nodes grouped for specific types of jobs. Each partition has its own limits such as maximum job runtime, hardware (CPUs, GPUs), and access permissions. Partitions are similar to queues in other workload managers. Programs (jobs) run via Slurm are sent to one or more of the physical servers (Nodes). Slurm is software that helps defining and executing these jobs, as well as managing users, permissions, and resource allocation. It helps track and display job details as well. Users must select a partition that matches the resources their job needs. If not specified, the default partition will be used.

**Slurm login servers** – the servers you connect to in order to submit your batch jobs.  
**Compute nodes** – the servers within the partitions where your jobs actually run.

### 1. BIU Cluster partitions Summary

#### Cluster General Partitions Summary

Partition Name	Purpose	Nodes	Max Time	Access (Accounts)	MaxJobs Per User	Max GPU /CPU Per USER	Node Mem
generic	General GPU jobs	dsicsgpu[02-09], dsisarit[02,05]	4h	All users	8 jobs	16 GPUs	192G
generic-48G	General GPU jobs	dsiaw15	4h	All users	2 jobs	2 GPUs	128G
H200-4h	GPU jobs on H200	hpc8h200-01	4h	All users	2 jobs	2 GPUs	2TB
H200-12h	GPU jobs on H200	hpc8h200-01	12h	All users	2 jobs	2 GPUs	2TB
L4-4h	GPU jobs on L4	hpc8l4-01-01	4h	All users	2 jobs	2 GPUs	512G
L4-12h	GPU jobs on L4	hpc8l4-01-01	12h	All users	2 jobs	2 GPUs	512G
A100-4h	GPU jobs on A100	hpc2a100-01, dsiasaf01	4h	All users	2 jobs	2 GPUs	512G
cpu1T-24h	CPU jobs	hpccpu01	24h	All Users	8 jobs	48 CPUs	1TB
cpu192G-48h	CPU jobs	dml[02-25]	48h	All Users	18 jobs	72 CPUs	192G
cpu512G-48h	CPU jobs	skittles, skittles[01-22]	48h	All Users	18 jobs	112 CPUs	512G
B200-4h	GPU jobs on B200	dgx-b200-01 dgx-b200-02	4h	All users	1 Job	2 GPUs	2TB
B200-12h	GPU jobs on B200	dgx-b200-01 dgx-b200-02	12h	All users	1 Job	2 GPUs	2TB
RTX6000-4h	GPU jobs on RTX6000	mm-lab02	4h	All users	2 Jobs	2 GPUs	2TB

**Cluster Private Partitions Summary**

Partition Name	Purpose	Nodes	Max Time	Access (Accounts)	MaxJobs Per User	Max GPU /CPU Per Job	Node Mem
p_kugler	GPU jobs on A100 Kugler group	hpc2a100-01	Unlimited	Users in Prof. <b>Kugler's</b> group	Unlimited	Unlimited	512G
p_amsterdamer	CPU jobs Amsterdamer group	dml[02-21]	Unlimited	Users in Prof. <b>Amsterdamer's</b> group	Unlimited	Unlimited	192G
p_glickman	GPU jobs Glickman group	dsicgpu[02–09]	Unlimited	Users in Prof. <b>Glickman's</b> group	Unlimited	Unlimited	192G
p_zohar	CPU jobs zohar group	skittles, skittles[01-22]	Unlimited	Users in Prof. <b>Zohar's</b> group	Unlimited	Unlimited	512G
p_kraus	GPU jobs Kraus group	dsisarit[02,05]	Unlimited	Users in Prof. <b>Kraus's</b> group	Unlimited	Unlimited	192G
p_weiss	GPU jobs Weiss group	dsiaw15	Unlimited	Users in Prof. <b>Weiss's</b> group	Unlimited	Unlimited	128G
p_peer	GPU jobs Peer group	dsiasaf01	Unlimited	Users in Prof. <b>Peer's</b> group	Unlimited	Unlimited	512G
p_rtx_schwartz	GPU jobs Schwartz group	mm-lab02	Unlimited	Users in Prof. <b>Schwartz's</b> group	Unlimited	Unlimited	2TB

Partition Name	Purpose	Nodes	Max Time	Access (Accounts)	MaxJobs MaxGPU Per User	Max GPU per Group	Node Mem
p_b200_chechik	GPU jobs on B200 Chechik group	dgx-b200-01	4h	Users in Prof. <b>Chechik's</b> group		1 GPU	2TB
p_b200_major	GPU jobs on B200 Major group	dgx-b200-01	4h	Users in Prof. <b>Major's</b> group		1 GPU	2TB
p_b200_kraus	GPU jobs on B200 Kraus group	dgx-b200-01	8h	Users in Prof. <b>Kraus's</b> group		1 GPU	2TB
p_b200_eng	GPU jobs on B200 Engineering group	dgx-b200-01	4h	Users in Prof. <b>Gannot, Fetaya, Lindenbaum, Tirer, Shavit</b> groups		1 GPU	2TB
p_b200_tsarfaty	GPU jobs on B200 Tsarfaty group	dgx-b200-01	4h	Users in Prof. <b>Tsarfaty's</b> group		1 GPU	2TB
p_b200_dagan	GPU jobs on B200 Dagan group	dgx-b200-01	4h	Users in Prof. <b>Dagan's</b> group		1 GPU	2TB
p_b200_goldberg	GPU jobs on B200 Goldberg group	dgx-b200-01	4h	Users in Prof. <b>Goldberg's</b> group		1 GPU	2TB
p_b200_nlp	GPU jobs on B200 NLP group	dgx-b200-01	4h	Users in Prof. <b>Tsarfaty, Dagan, Goldberg</b> groups		4 GPU	2TB
p_b200_schwartz	GPU jobs Schwartz group	dgx-b200-02	Unlimited	Users in Prof. <b>Schwartz's</b> group	Unlimited	Unlimited	2TB

**Note:** These partitions are restricted. Only users who are members of the matching group can submit jobs to that partition.

**Scheduling and Limits:** Jobs submitted to these private partitions receive higher scheduling priority on their dedicated resources. Other limits are listed in the table in Section 1.

For more details, See Section 6 in this document - **Private Partitions and Accounts – How to submit jobs**

## 2. Job Submission Examples

**Note:** All job submission scripts must include lines that begin with #SBATCH. These are special directives that Slurm uses to allocate resources and control job behavior.

This section explains how to create a slurm job script and submit it, including examples and explanations of the parameters.

### Basic CPU Job

```
#!/bin/bash
#SBATCH --job-name=cpu_test
#SBATCH --output=cpu_test_%j.out
#SBATCH --error=cpu_test_%j.err
#SBATCH --partition=<Partition-name>
#SBATCH --cpus-per-task=4
#SBATCH --mem=<size>[M- G]
< user commands>
< example: python script.py>
```

### Basic GPU Job

```
#!/bin/bash
#SBATCH --job-name=gpu_test
#SBATCH --output=gpu_test_%j.out
#SBATCH --error=gpu_test_%j.err
#SBATCH --partition=<partition name>
#SBATCH --gres=gpu:1
#SBATCH --mem=<size>[M G]
<user commands>
<example: python gpu_script.py>
```

### Resources:

- By default, **1 CPU** is allocated to a job. To request more CPUs use:  
**#SBATCH --cpus-per-task=N**, where *N* is the number of CPUs you need.  
Ensure that *N* does not exceed the CPU limit of the partition, which can be found in the partition table in Section 1.
- By default, 16 GB of memory is allocated for each CPU assigned to the job.  
To request more memory use: **#SBATCH --mem=<size>[M- G]**
- By default, **no GPUs** are allocated to a job. To request GPUs use:  
**#SBATCH --gres=gpu:1**, where *N* is the number of GPUs you need.  
Ensure that *N* does not exceed the GPU limit of the partition, which can be found in the partition table in Section 1.

in case the partition includes several types of GPU's and you want a specific GPU, use:  
**#SBATCH --gres=gpu:<type>:<count>**

### Email notification:

You can declare (enable) email notifications in your job Slurm script file:

```
#SBATCH --mail-user=your.email@example.com  
#SBATCH --mail-type=[ALL,BEGIN,END,FAIL]
```

### Maximum run time:

You can use the --time option to set the maximum runtime for your job.

```
#SBATCH --time=HH:MM:SS
```

- If you don't specify --time, the job will run using the **default time limit of the partition**.
- In case you use --time option, please make sure the time you request **does not exceed** the maximum allowed by your partition.

### Docker Job

If you are using docker:

- Run the container in the foreground (**without using -d**).
- Docker container name should be **slurm-\$SLURM\_JOB\_ID** for traceability  
Use this in your batch script:  
**DockerName=slurm-job-\$SLURM\_JOB\_ID**  
docker run --name "\$DockerName" --rm my-image:latest python script.py

- If your Slurm job requests **GPUs**, use the list that Slurm provides in **SLURM\_JOB\_GPUS** variable to restrict Docker to the GPUs allocated to you.

#### In your sbatch job script:

```
docker run --name "$DockerName" --rm --gpus "device=${SLURM_JOB_GPUS}"  
my-image:latest python script.py
```

SLURM\_JOB\_GPUS is typically a comma-separated list of device indices (e.g., 0,2).

- If your workflow relies on **CUDA\_VISIBLE\_DEVICES**, pass it into the container: CUDA\_VISIBLE\_DEVICES="\${SLURM\_JOB\_GPUS}"

### Example of a Docker batch script:

```
#!/bin/bash  
#SBATCH --job-name=docker_job  
#SBATCH --output=docker_job-%j.out  
#SBATCH --error=docker_job-%j.err  
#SBATCH --partition=generic  
#SBATCH --gres=gpu:1  
DockerName=slurm-job-$SLURM_JOB_ID  
docker run --name "$DockerName" --rm --gpus "device=${SLURM_JOB_GPUS}" \  
my-image:latest python script.py
```

**Conda Job**

```
#!/bin/bash
#SBATCH --job-name=conda_job
#SBATCH --output=conda_job_%j.out
#SBATCH --error=conda-job_%j.err
#SBATCH --partition=<partition name>
#SBATCH --mem=<size>[M G]
source ~/miniconda3/etc/profile.d/conda.sh
conda activate myenv # Or source your environment
python script.py
```

**MATLAB job**

```
#!/bin/bash
#SBATCH --job-name=matlab_job
#SBATCH --output=matlab_job_%j.out
#SBATCH --error=matlab-job_%j.err
#SBATCH --partition=<partition name>
#SBATCH --mem=<size>[M G]
#SBATCH --cpus-per-task=48
<user commands>
matlab -batch "run('mymatlab.m')"
<user commands>
```

**Important Note:**

When Running MATLAB with **parallel workers** and submitting a MATLAB job with Slurm using:

**#SBATCH --cpus-per-task=48**      <-- 48 is just an example

In order to ensure MATLAB uses **all the CPU cores allocated by Slurm**, add the following to your MATLAB job script (e.g. **mymatlab.m**):

```
num_workers = str2double(getenv('SLURM_CPUS_PER_TASK'));
c = parcluster('local');
c.NumWorkers = num_workers;    % or exact number of cores you want to use
(e.g. 48)
saveProfile(c);
```

### 3. How to submit the job

1. Connect to one of the slurm\_login servers suing SSH:  
`ssh slurm-login1.lnx.biu.ac.il` or  
`ssh slurm-login2.lnx.biu.ac.il` or  
`ssh slurm-login3.lnx.biu.ac.il`
2. Submit Your Job: **sbatch myJob.sh**
3. Monitor your job using **squeue -u <your-username>**
4. After submission, Slurm will generate output and error files according to what you have declared in your sbatch job:

```
#SBATCH --output=<file-name>.out  
#SBATCH --error=<file-name>.err
```

Example:

```
#SBATCH --output=matlab_job_%j.out # the file name include the job-id  
#SBATCH --error=matlab-job_%j.err # the file name include the job-id
```

**Make sure you check output and error files for progress**

### 4. Job Time Limits and Requeue Policy

Jobs submitted with **sbatch** will be automatically **suspended and requeued** by the system when they reach the time limit of the partition they were submitted to (see the Max time in Part 1).

#### Checkpoints:

To benefit from requeuing, implement **checkpointing** in your application:

- Save progress periodically.
- On restart, resume from the latest checkpoint.

**Without checkpointing, the requeued job will start from the beginning.**

## 5. Interactive Jobs

srun is used to run immediate/interactive commands directly on compute nodes managed by Slurm.

srun can both request resources and immediately run a command in that allocation.

### Examples:

1. **srun --partition=<partition-name> <command>**  
Runs a single command on a compute node in the specified partition
2. **srun --partition=<partition-name> --pty bash**  
Starts an **interactive shell** on a compute node in the given partition.  
When you exit the shell (exit), the allocation ends and resources are released.
3. **srun --partition=<partition-name> --gres=gpu:1 --time=HH:MM:SS --pty bash**  
Open an interactive shell with 1 GPU for 1 hour.  
When you exit the shell (exit), the allocation ends and resources are released.

### Important Notes:

- Interactive jobs launched this way will **not** be automatically requeued.
- Use **srun** only for quick tests, debugging, or development. For longer jobs or those requiring resilience to timeouts, use **sbatch** with checkpointing

## 6. Private Partitions and Accounts – How to submit jobs

The table below lists the private partitions.

If your group appears in the “**Group**” column, you can use the corresponding partition in the first column, “**Partition**”, with the account shown in the “**Account**” column.

The exact sbatch command is provided in the “SBATCH commands” column.”

To check which Slurm accounts (groups) you belong to, run:

**sacctmgr list assoc where user=\$USER format=User,Account -nP**

Replace \$USER with your actual username.

Users who are not part of the appropriate group **cannot** use these private partitions.

Include the matching partition and account in your batch script:

**#SBATCH --partition=<partition-name>**

**#SBATCH --account=<account-name>**

For the exact values, refer to the “SBATCH commands” column in the table below.

Partition	Group	Account	SBATCH commands
p_kugler	Kugler group	ug_kugler	#SBATCH --partition=p_kugler #SBATCH --account=ug_kugler
p_amsterdamer	Amsterdamer group	ug_amsterdamer	#SBATCH --partition=p_amsterdamer #SBATCH --account=ug_amsterdamer
p_glickman	Glickman group	ug_glickman	#SBATCH --partition=p_glickman #SBATCH --account=ug_glickman
p_zohar	Yoni Zohar group	ug_zohar	#SBATCH --partition=p_zohar #SBATCH --account=ug_zohar
p_kraus	Kraus group	ug_kraus	#SBATCH --partition=p_kraus #SBATCH --account=ug_kraus
p_weiss	Weiss group	ug_weiss	#SBATCH --partition=p_weiss #SBATCH --account=ug_weiss
p_peer	Peer group	ug_peer	#SBATCH --partition=p_peer #SBATCH --account=ug_peer
p_rtx_schwartz	Schwartz group	ug_schwartz	#SBATCH --partition=p_rtx_schwartz #SBATCH --account=ug_schwartz
p_b200_chechik	Chechik	ug_chechik	#SBATCH --partition=p_b200_chechik #SBATCH --account=ug_chechik
p_b200_major	Major	ug_major	#SBATCH --partition=p_b200_major #SBATCH --account=ug_major
p_b200_kraus	Kraus	ug_kraus	#SBATCH --partition=p_b200_kraus #SBATCH --account=ug_kraus
p_b200_eng	Gannot group Fetaya group Tirer group Shavit group Lindenbaum group	ug_gannot ug_fetaya ug_Tirer ug_shavit ug_lindenbaum	#SBATCH --partition=p_b200_eng #SBATCH --account=ug_gannot #SBATCH --account=ug_fetaya #SBATCH --account=ug_Tirer #SBATCH --account=ug_shavit #SBATCH --account=ug_lindenbaum
p_b200_tsarfaty	Tsarfaty	ug_tsarfaty	#SBATCH --partition=p_b200_tsarfaty #SBATCH --account=ug_tsarfaty
p_b200_dagan	Dagan	ug_dagan	#SBATCH --partition=p_b200_nlp #SBATCH --account=ug_dagan
p_b200_goldberg	Goldberg	ug_goldberg	#SBATCH --partition=p_b200_nlp #SBATCH --account=ug_goldberg
p_b200_nlp	Tsarfaty group Dagan group Goldberg group	ug_tsarfaty ug_dagan ug_goldberg	#SBATCH --partition=p_b200_nlp #SBATCH --account=ug_tsarfaty #SBATCH --account=ug_dagan #SBATCH --account=ug_goldberg
p_b200_schwartz	Schwartz group	ug_schwartz	#SBATCH --partition=p_b200_schwartz #SBATCH --account=ug_schwartz

## 7. Using Jupyter on Slurm Cluster

1. Install Jupyter Notebook locally in your home folder: **pip3 install notebook** (if it is not installed already).

2. Start Jupyter (no browser) on the compute node:

Connect to one of the **slurm\_login servers** and run the following command:

**srun --partition=<partition> --gres=gpu:1 jupyter-notebook --time=02:00:00 --no-browser --ip=0.0.0.0 &**

- Note the **node name** and **port number** assigned by Slurm.
- Copy and save the URL that Jupyter prints (you'll need it later)

**output example:**

`http://nodename:8888/tree?token=e34d73f70fa48254500e0556663db7b859985448f9d9829d`

`http://127.0.0.1:8888/tree?token=e34d73f70fa48254500e0556663db7b859985448f9d9829d`

In this example **node name** and **port number**, assigned by Slurm, are: nodename and 8888

3. On your local PC - Create an SSH tunnel:

On your local PC, run:

**ssh -L <port-num>:localhost:<port-num> -J <username>@<slurm\_login\_server> <username>@<node>**

- Replace <port-num> with the actual port number from step 1.
- Replace <username> with your username.
- Replace <node> with the node assigned by Slurm (from step 1).

**Use Jupyter Notebook from a Web Browser:**

paste the copied Jupyter URL (from step 1, second line) into your **local web browser**.

You should now see the notebook interface running on the remote Slurm node.

**Use Jupyter Notebook inside VS Code:**

On your local PC, open VS Code and install the Microsoft **Python** and **Jupyter** extensions (if they are not already installed). (No Remote-SSH is needed)

**To create a new notebook file:**

- Ctrl + Shift + P → Jupyter: Create New Jupyter Notebook
- On the top-right corner and click Select Kernel.
- Choose → Existing Jupyter Server...

Paste the remote Jupyter server URL the one that you copied when started jupyter)  
for example: `http://localhost:8888/?token=abcd1234... )`

VS Code will now connect to the Jupyter Notebook running on your Slurm node.

**Important Notes:**

- After testing and debugging your code, submit the **actual job** to the desired partition using **sbatch** and a Slurm **batch script**.
- It is recommended to perform **debugging and testing** in a partition with **less powerful CPU/GPU nodes**, to avoid occupying high-demand resources during development. For **production jobs**, use sbatch with sbatch scripts.
- Using `--time=HH:MM:SS` will make sure the Job will stop automatically when the `-time` limit is reached or the partition's time limit is hit.

## 8. Useful Slurm Commands

Submit a job	<b><i>sbatch job.sh</i></b>
Check your jobs	<b><i>squeue -u \$USER</i></b>
Cancel a job	<b><i>scancel &lt;job_id&gt;</i></b>
View available partitions	<b><i>sinfo</i></b>

---

Version 1.8, 3 Dec 2025